

# Performance Analysis of a Live Sharded Blockchain System

The logo for Multiversx, featuring the word "Multivers" in white and a stylized "x" in teal, all contained within a black rounded rectangle.

Multiversx

Debdoot Roy Chowdhury (23M0675)  
Guided By: Prof. Vinay J. Ribeiro

# What is MultiVersX?

- **MultiversX**, previously known as Elrond, **founded in 2019** by Benjamin Mincu.
- It is a **distributed blockchain network** for web3 applications that utilizes a **sharded state architecture** and secure **Proof of Stake** consensus mechanism.
- Their main network has over **3000 active validators** and over **400M confirmed transactions**.
- They provide efficient **ecosystem to create dApps** within 30 minutes, and has partnered with many popular **Google Cloud, Tencent, AWS, Digital Ocean, Coinbase** e.t.c.

# Terminologies and Architecture



- Time is divided into **fixed time slots** of Epochs and Rounds.
- **Each Epoch** consist of **multiple rounds** of 6 Seconds.

# Terminologies and Architecture



- Time is divided into **fixed time slots** of Epochs and Rounds.
- **Each Epoch** consist of **multiple rounds** of 6 Seconds.
- Each Shard produces **at most 1 Block in a Round**.


# Terminologies and Architecture



- Time is divided into **fixed time slots** of Epochs and Rounds.
- **Each Epoch** consist of **multiple rounds** of 6 Seconds.
- Each Shard produces **at most 1 Block in a Round**.
  - 100 Rounds compiles to 1 Epoch.
- After each Epoch, **shard reorganization and pruning** takes place.

# Terminologies and Architecture


→ Each **Block** consists of several **Mini-Blocks**.



Block

Mini  
Blocks

# Terminologies and Architecture




Block

Mini  
Blocks

- Each **Block** consists of several **Mini-Blocks**.
- Each **Mini-Block** consists of similar type of **Transactions** (Txns to a specific shard is stored in one MiniBlock).

# Terminologies and Architecture




Block

Mini  
Blocks

- Each **Block consists of several Mini-Blocks**.
- Each **Mini-Block consists of similar type of Transactions** (Txns to a specific shard is stored in one MiniBlock).
- Each Mini-Block can hold at most 5000 transactions. And each block can hold at most 20000 transactions.




# Terminologies and Architecture



- Each **Block consists of several Mini-Blocks**.
- Each **Mini-Block consists of similar type of Transactions** (Txns to a specific shard is stored in one MiniBlock).
- Each Mini-Block can hold at most 5000 transactions. And each block can hold at most 20000 transactions.
- It is useful to **segregate different transactions into different mini-blocks** and speed up fetching and storing.

# Terminologies and Architecture


Belongs to Shard k



Mini-Block  
consisting of  
transactions  
from  
Shard k to n


# Terminologies and Architecture

## Shard 0




# Terminologies and Architecture

## Shard 0




## Shard 1




# Terminologies and Architecture


## Shard 0



## Shard 1



## Meta Shard



# Terminologies and Architecture

→ “Secure” Proof of Stake is used for Consensus.

# Terminologies and Architecture

- “Secure” Proof of Stake is used for Consensus.
- Adaptive State Sharding technique is followed to dynamically change the number of shards.

# Terminologies and Architecture


- “Secure” Proof of Stake is used for Consensus.
- Adaptive State Sharding technique is followed to dynamically change the number of shards.
- Follows Account-Based transaction model.



# Terminologies and Architecture


- “Secure” Proof of Stake is used for Consensus.
- Adaptive State Sharding technique is followed to dynamically change the number of shards.
- Follows Account-Based transaction model.
- Each shard process transactions of a distinct set of accounts based on the “bech32” address provided to that account.

# Adaptive State Sharding




Starts with  
1 Shard

# Adaptive State Sharding





Starts with  
1 Shard




Threshold  
Crossed

# Adaptive State Sharding




Threshold  
Crossed



$S_0$  divided to two shards  
 $S_0$  and  $S_1$


# Adaptive State Sharding



Starts with  
1 Shard



Threshold  
Crossed



$S_0$  divided to two shards  
 $S_0$  and  $S_1$

Branch encoding from leaf to root, represents the last n-bits of the account addresses that will have their originating transaction processed by that shard.

# Adaptive State Sharding

Addresses



Mapped to  $S_0$

# Adaptive State Sharding

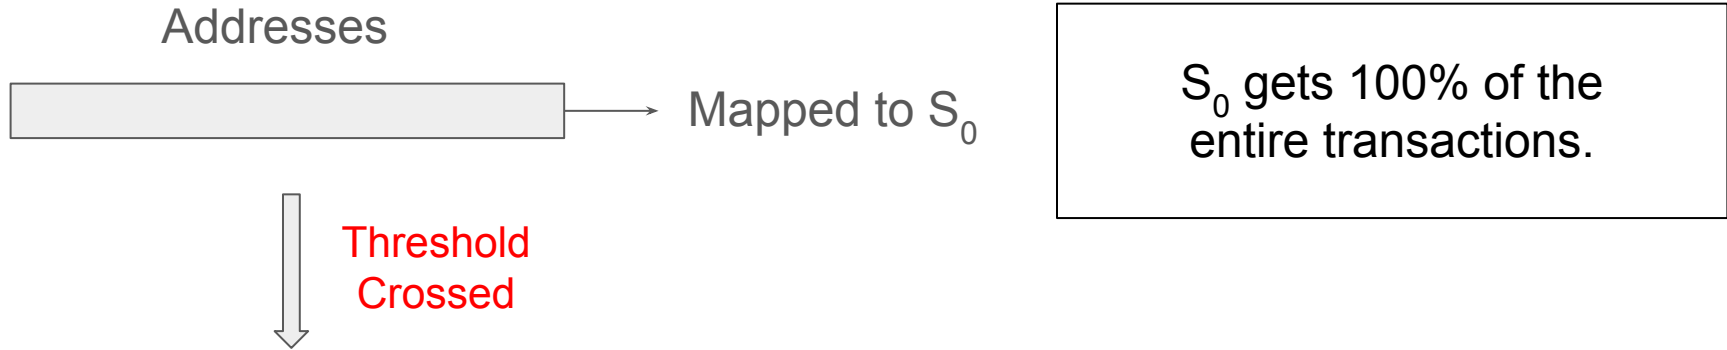
Addresses



Mapped to  $S_0$


$S_0$  gets 100% of the  
entire transactions.

# Adaptive State Sharding





# Adaptive State Sharding



$S_0$  gets 100% of the entire transactions.

# Adaptive State Sharding

Addresses



→ Mapped to  $S_0$

$S_0$  gets 100% of the  
entire transactions.



Threshold  
Crossed



0

→ Mapped to  $S_0$




1



→ Mapped to  $S_1$

$S_0$  and  $S_1$  each gets 50%  
of the entire transactions.

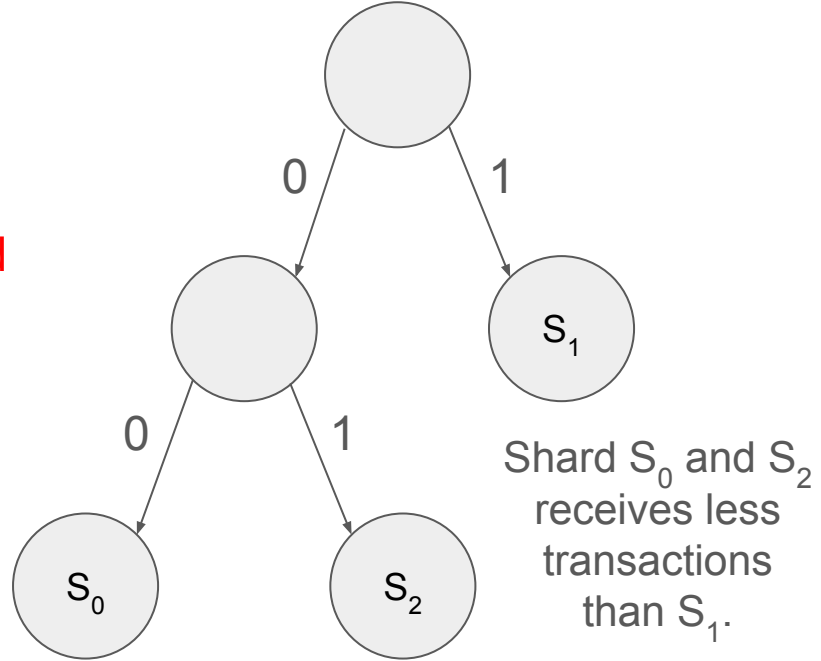
# Adaptive State Sharding



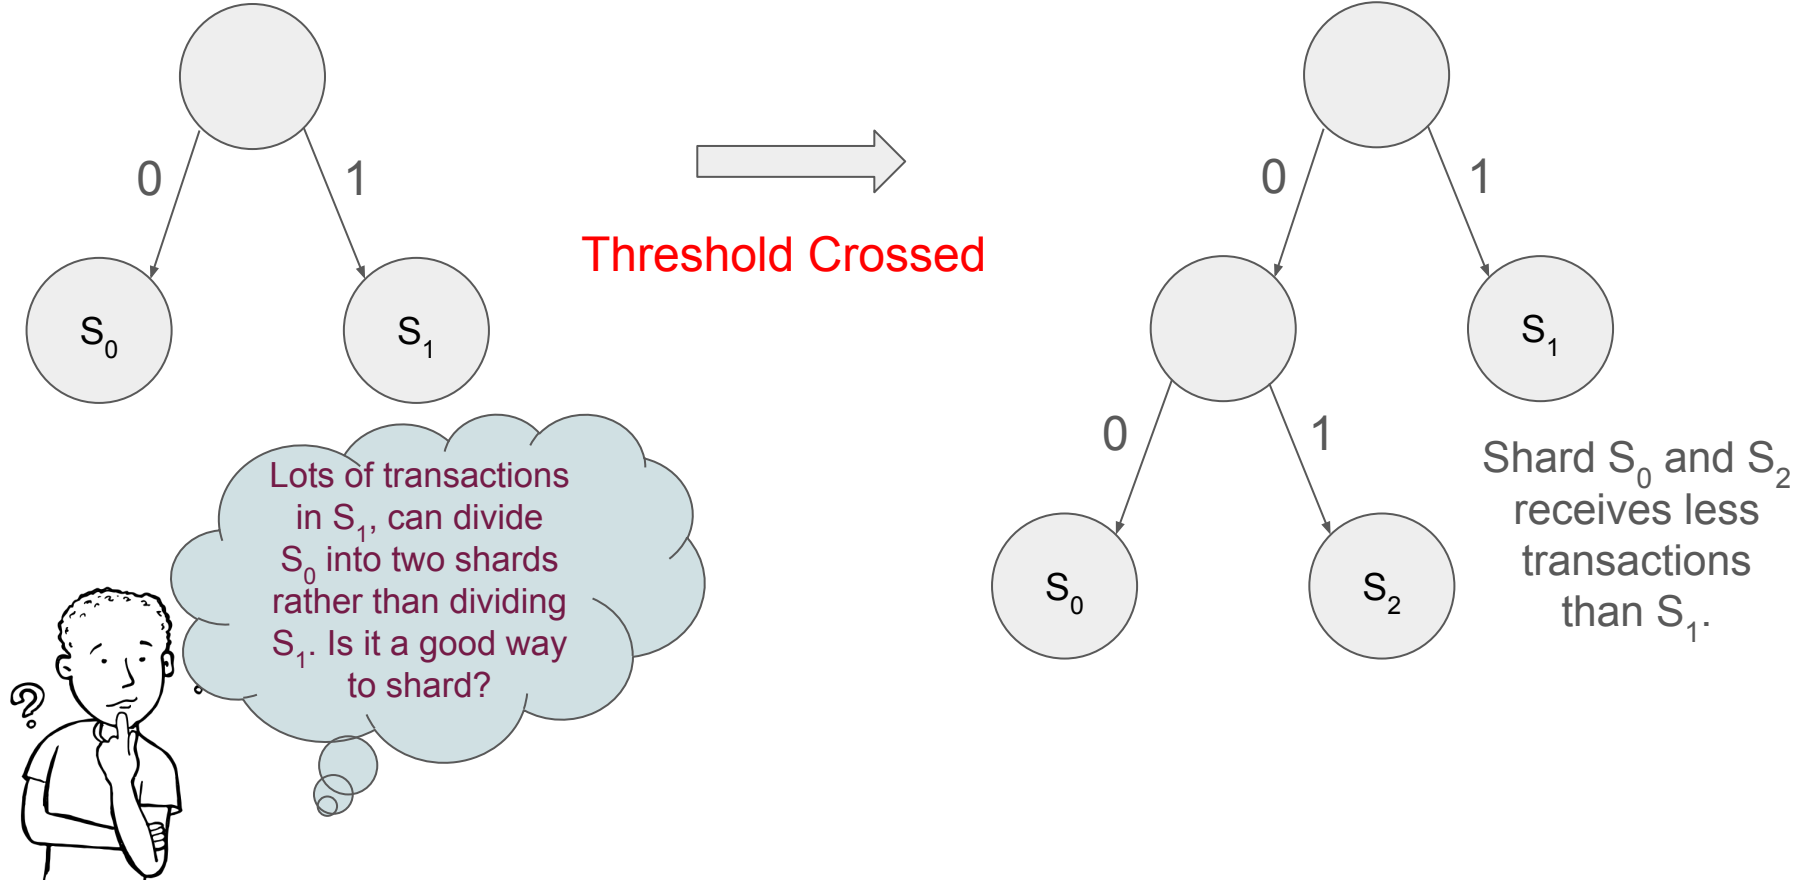
# Adaptive State Sharding



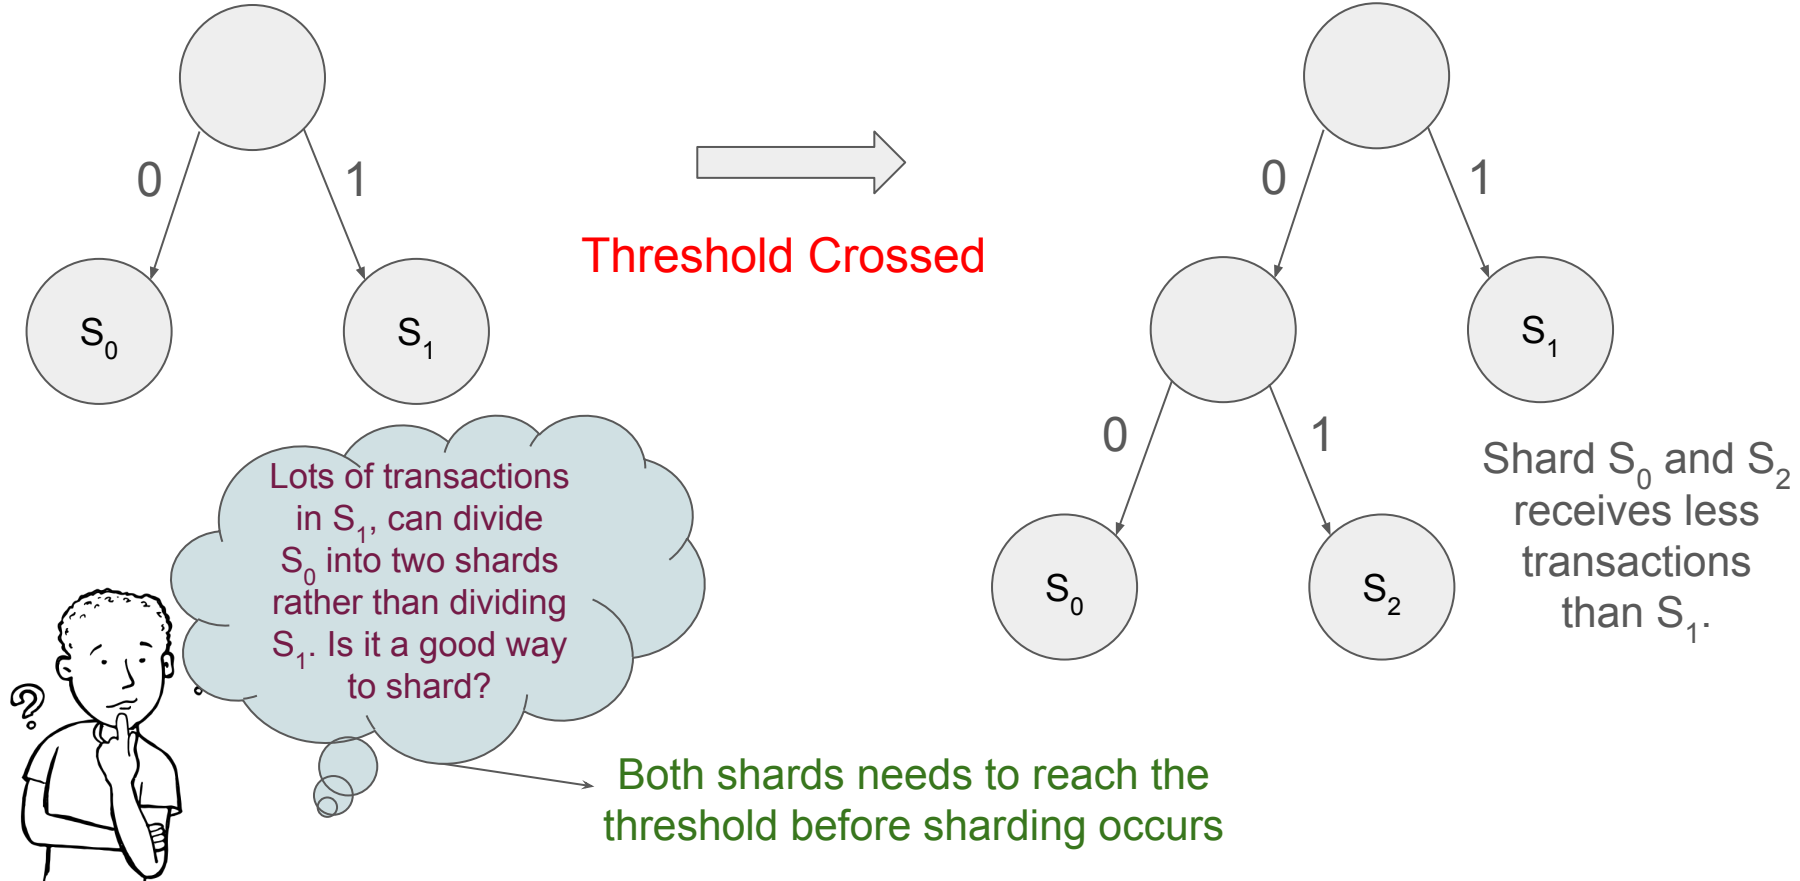
Threshold Crossed



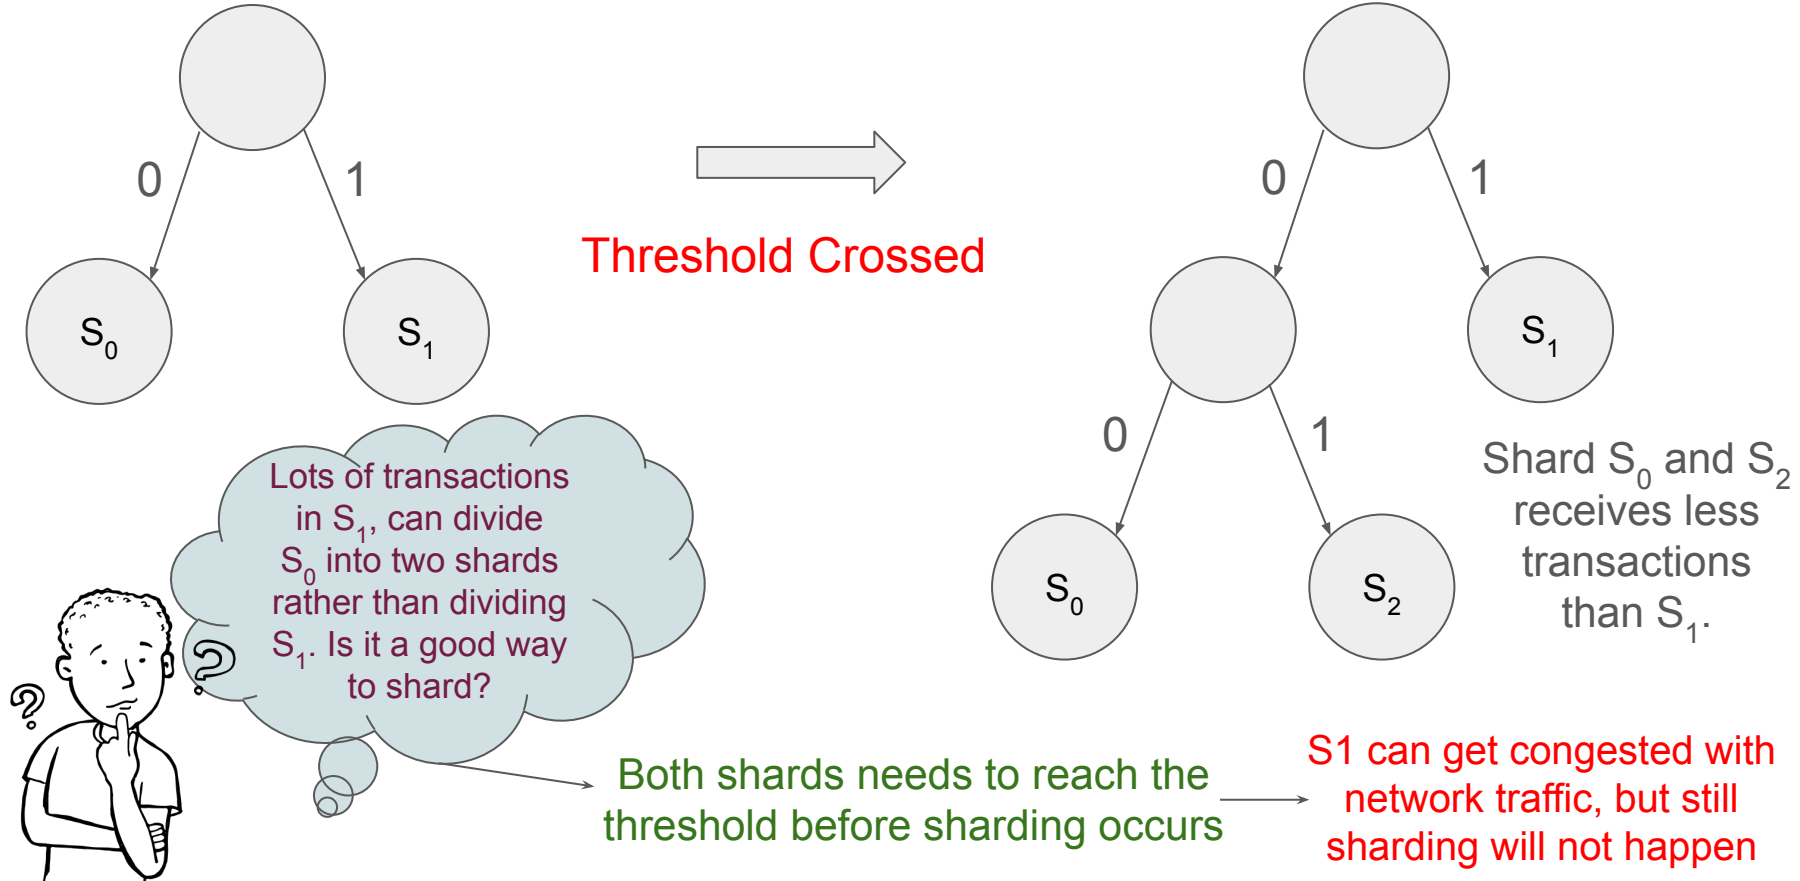
# Adaptive State Sharding



# Adaptive State Sharding




# Adaptive State Sharding



# Adaptive State Sharding


Addresses





# Adaptive State Sharding


Addresses




$S_0$  and  $S_1$  each gets 50% of the entire transactions.

# Adaptive State Sharding

Addresses




Threshold  
Crossed




$S_0$  and  $S_1$  each gets 50%  
of the entire transactions.

# Adaptive State Sharding

Addresses



$S_0$  and  $S_1$  each gets 50% of the entire transactions.





$S_0$  and  $S_2$  each gets 25% of the entire transactions, while  $S_1$  gets the rest 50%.

# Secure Proof of Stake

- Random Validators' Selection + Eligibility through Stake + Rating + Optimal Dimension for the Consensus Group.


# Secure Proof of Stake

- Random Validators' Selection + Eligibility through Stake + Rating + Optimal Dimension for the Consensus Group.
- For each shard, there will be two lists of Nodes - Eligible List and Waiting List.




# Secure Proof of Stake

- Random Validators' Selection + Eligibility through Stake + Rating + Optimal Dimension for the Consensus Group.
- For each shard, there will be two lists of Nodes - Eligible List and Waiting List.



# Secure Proof of Stake

- Random Validators' Selection + Eligibility through Stake + Rating + Optimal Dimension for the Consensus Group.
- For each shard, there will be two lists of Nodes - Eligible List and Waiting List.



# Secure Proof of Stake


→ How does the network supports such high throughput?

- Using small consensus groups to reduce communication overhead.
- Keeping transaction size less.
- Shard pruning keeps the nodes light.




# Secure Proof of Stake

Node probability (size)  
based on **stake** and **rating**



# Secure Proof of Stake


Node probability (size)  
based on **stake** and **rating**



# Secure Proof of Stake

Node probability (size)  
based on **stake** and **rating**

Consensus  
group  
selection  
function




Last block's signature

# Secure Proof of Stake

Node probability (size)  
based on **stake** and **rating**

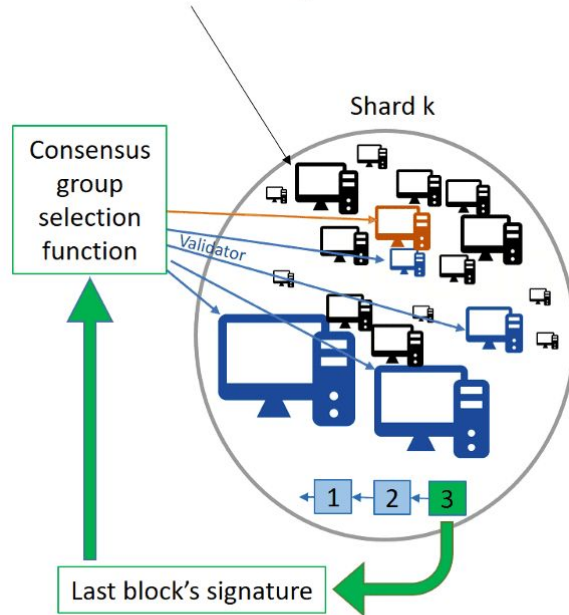
Consensus  
group  
selection  
function



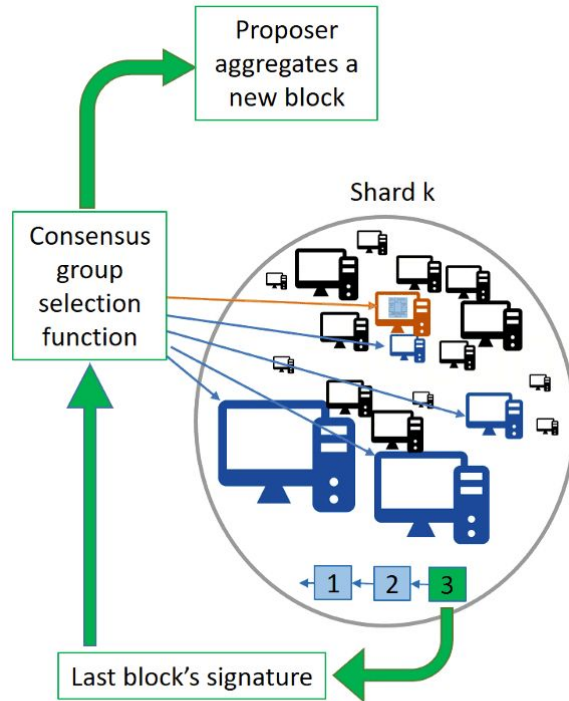
Last block's signature

# Secure Proof of Stake

Node probability (size)  
based on **stake** and **rating**

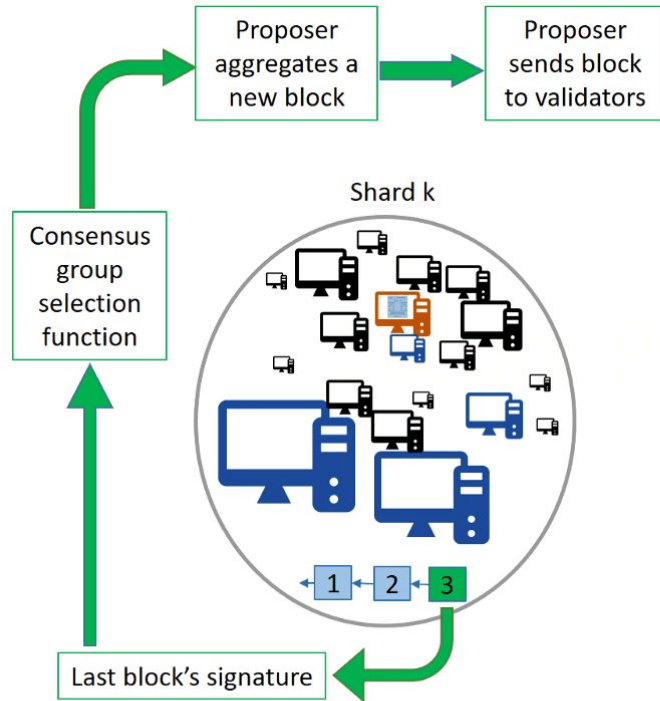


# Secure Proof of Stake



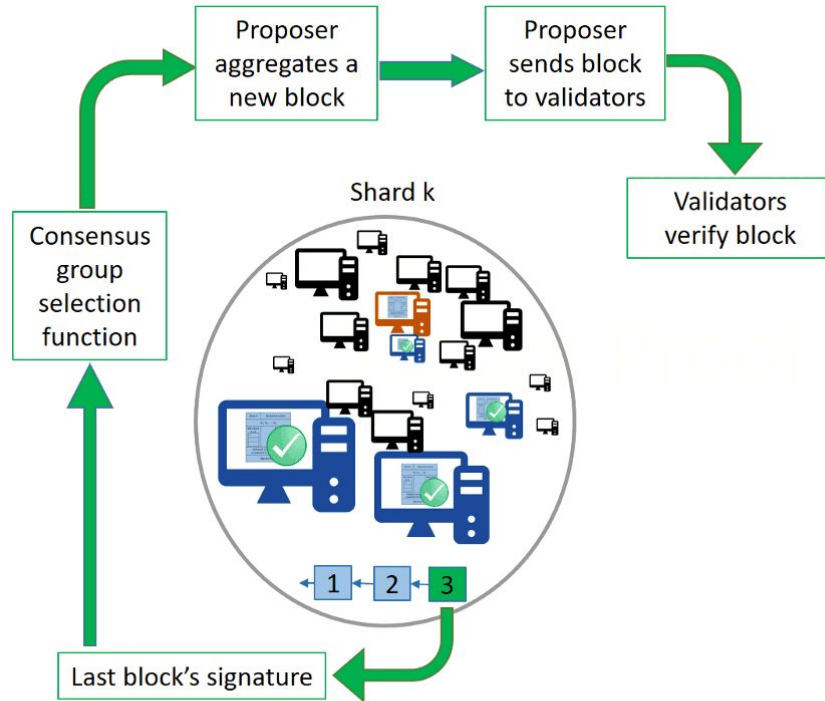
Block 4	Header	
$Tx_1, Tx_2, \dots, Tx_n$		
Mini Block $k \rightarrow 1$	...	Mini Block $k \rightarrow n$
<div>Tx</div>		<div>Tx</div>
<div></div>		<div></div>
<div></div>		<div></div>
Validated Tx Inclusion Proofs Invalidated Tx Inclusion Proofs		
Signature: null		

# Secure Proof of Stake



Block 4	Header	
$Tx_1, Tx_2, \dots, Tx_n$		
Mini Block $k \rightarrow 1$	...	Mini Block $k \rightarrow n$
<div>Tx</div>		<div>Tx</div>
<div></div>		<div></div>
<div></div>		<div></div>
Validated Tx Inclusion Proofs		
Invalidated Tx Inclusion Proofs		
Signature: null		

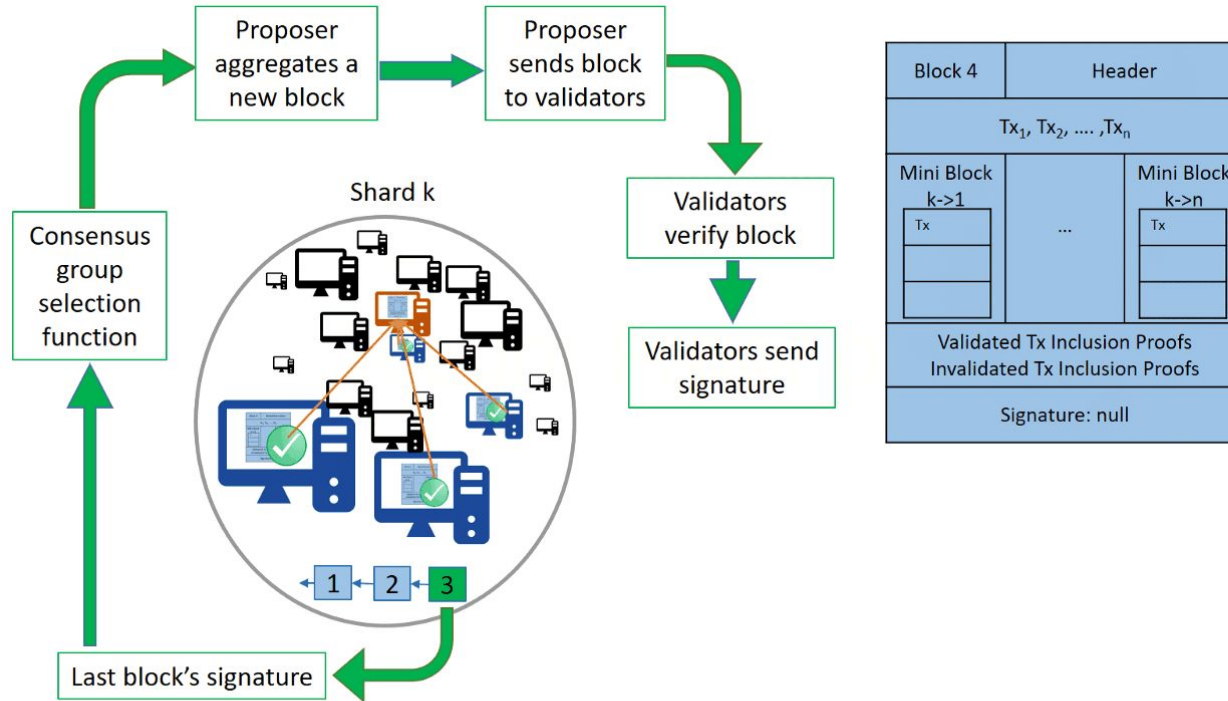
# Secure Proof of Stake




Block 4	Header	
$Tx_1, Tx_2, \dots, Tx_n$		
Mini Block $k \rightarrow 1$	...	Mini Block $k \rightarrow n$
<div>Tx</div>		<div>Tx</div>
<div></div>		<div></div>
<div></div>		<div></div>
Validated Tx Inclusion Proofs		
Invalidated Tx Inclusion Proofs		
Signature: null		



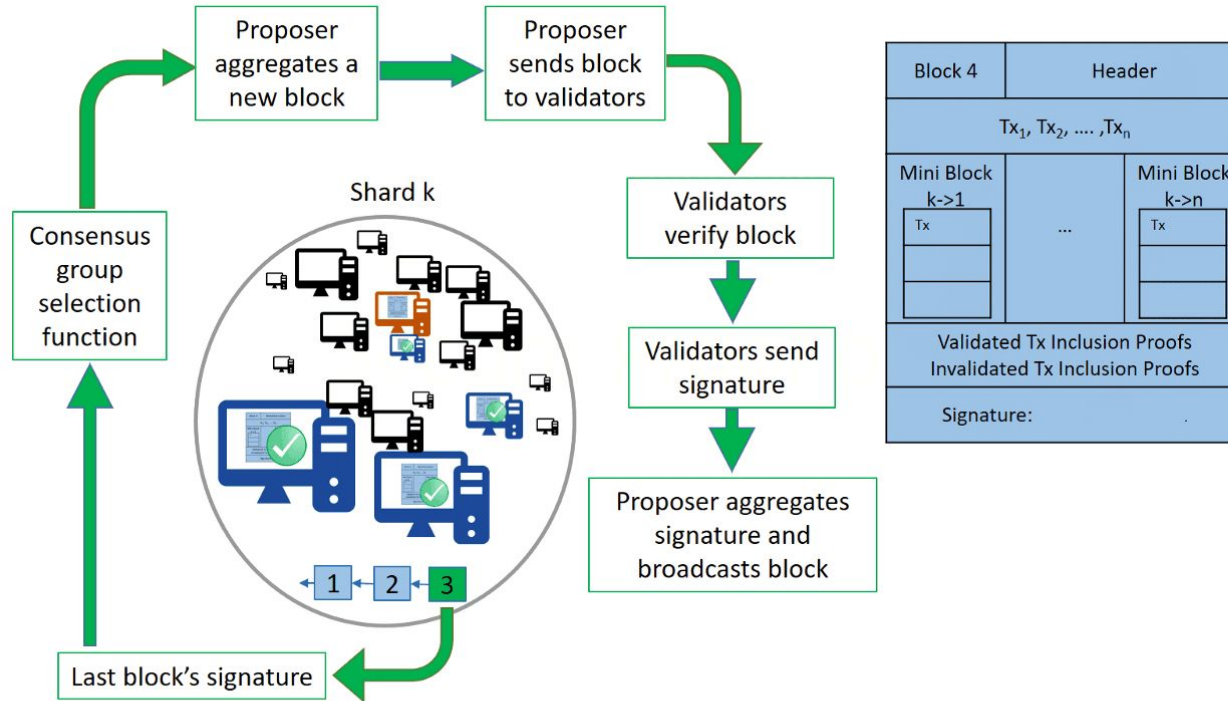
# Secure Proof of Stake



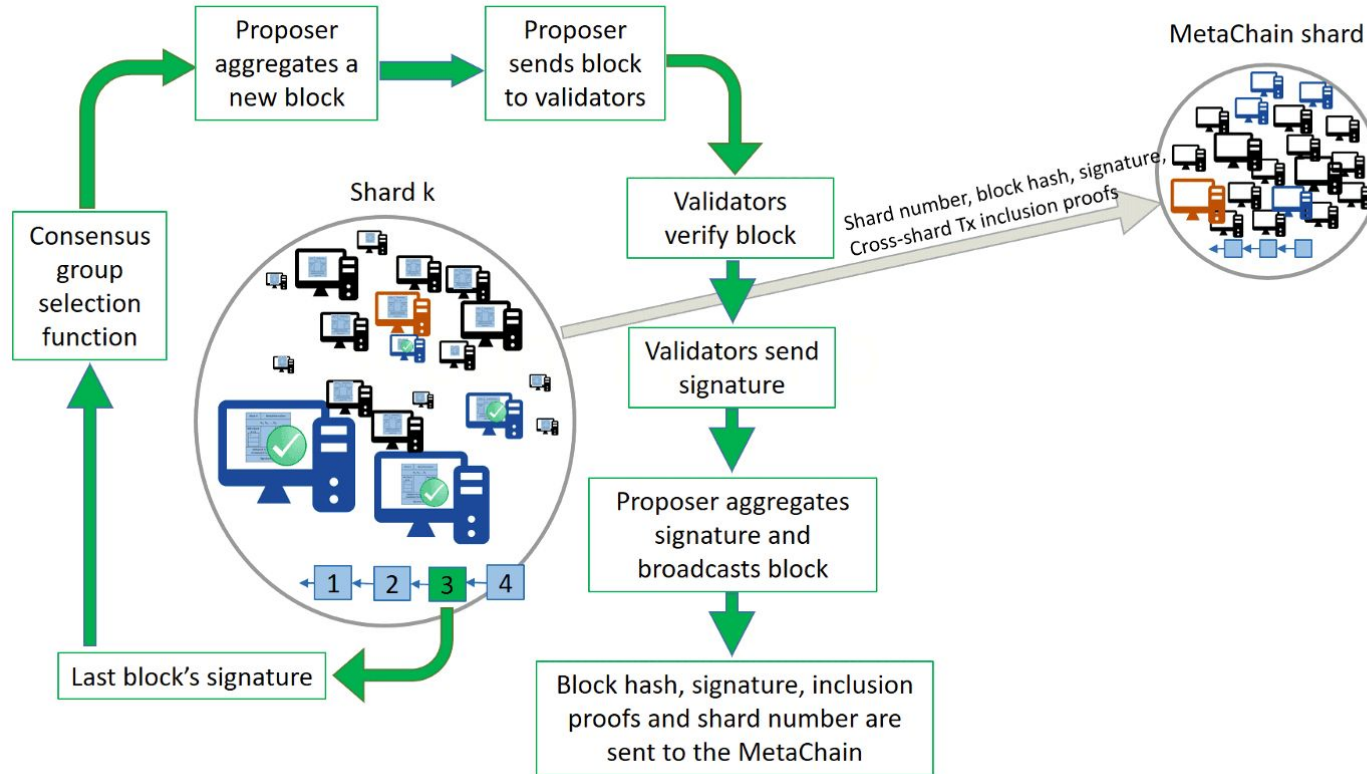
# Secure Proof of Stake



# Secure Proof of Stake



# Secure Proof of Stake




# Security




block m

# Security

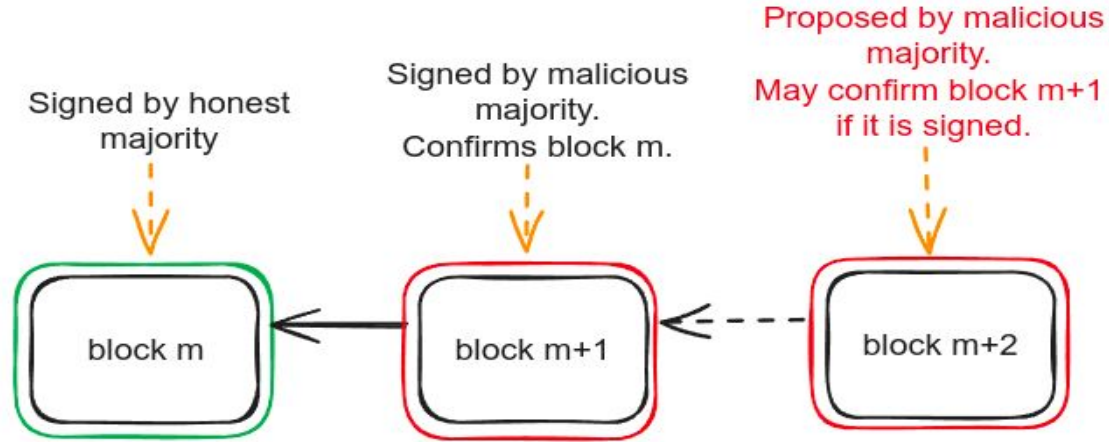
Signed by honest  
majority



# Security

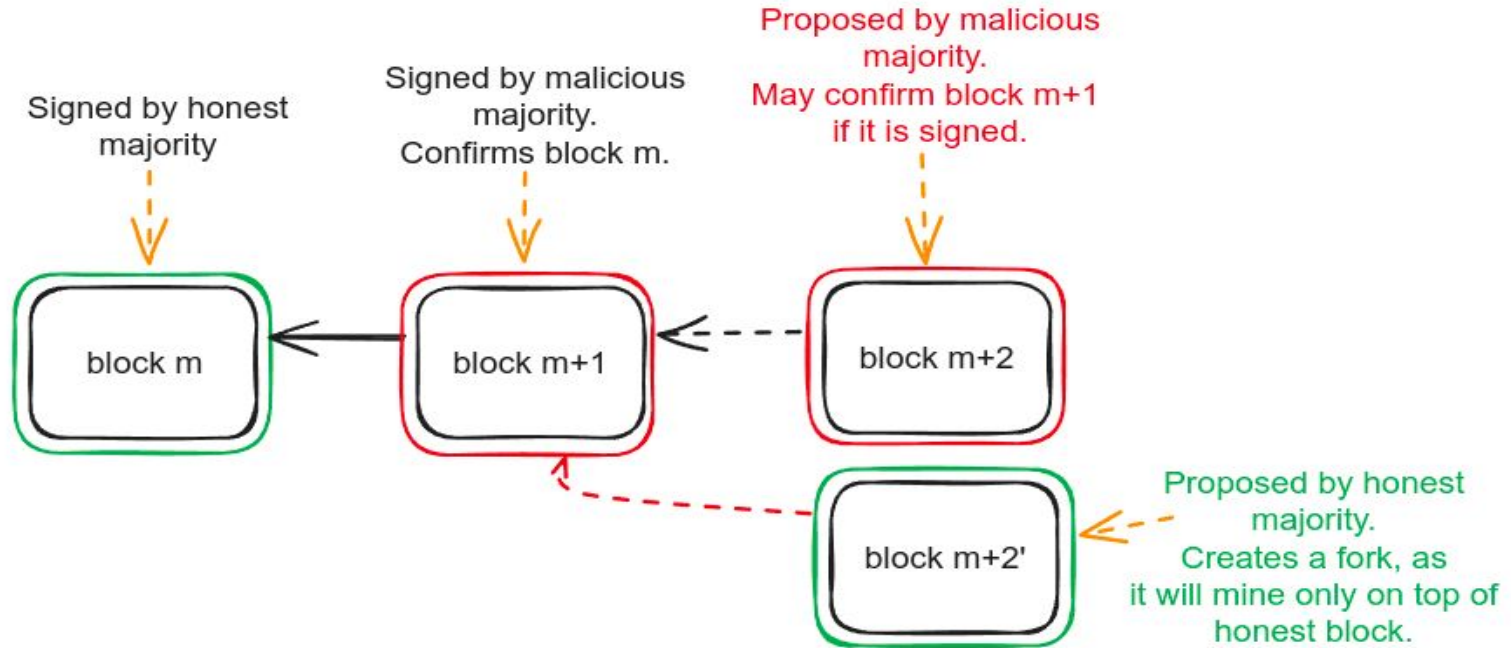


# Security

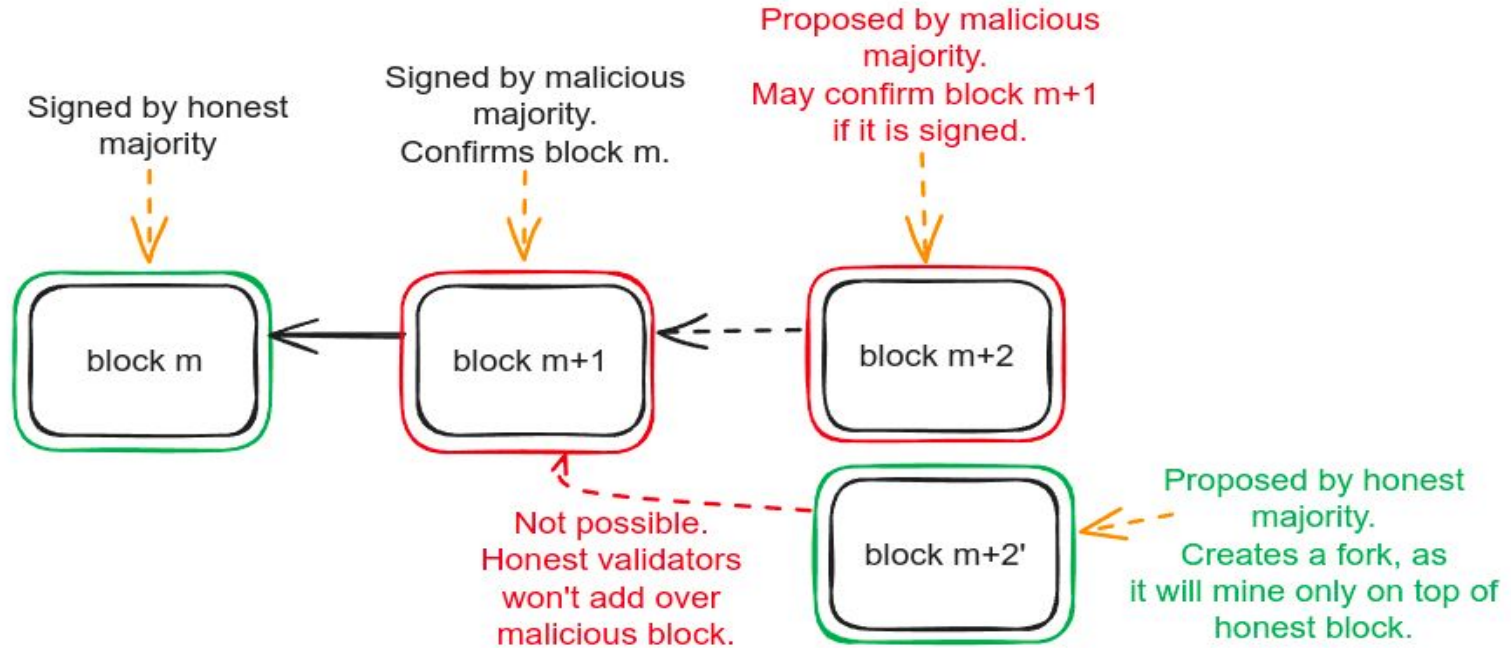




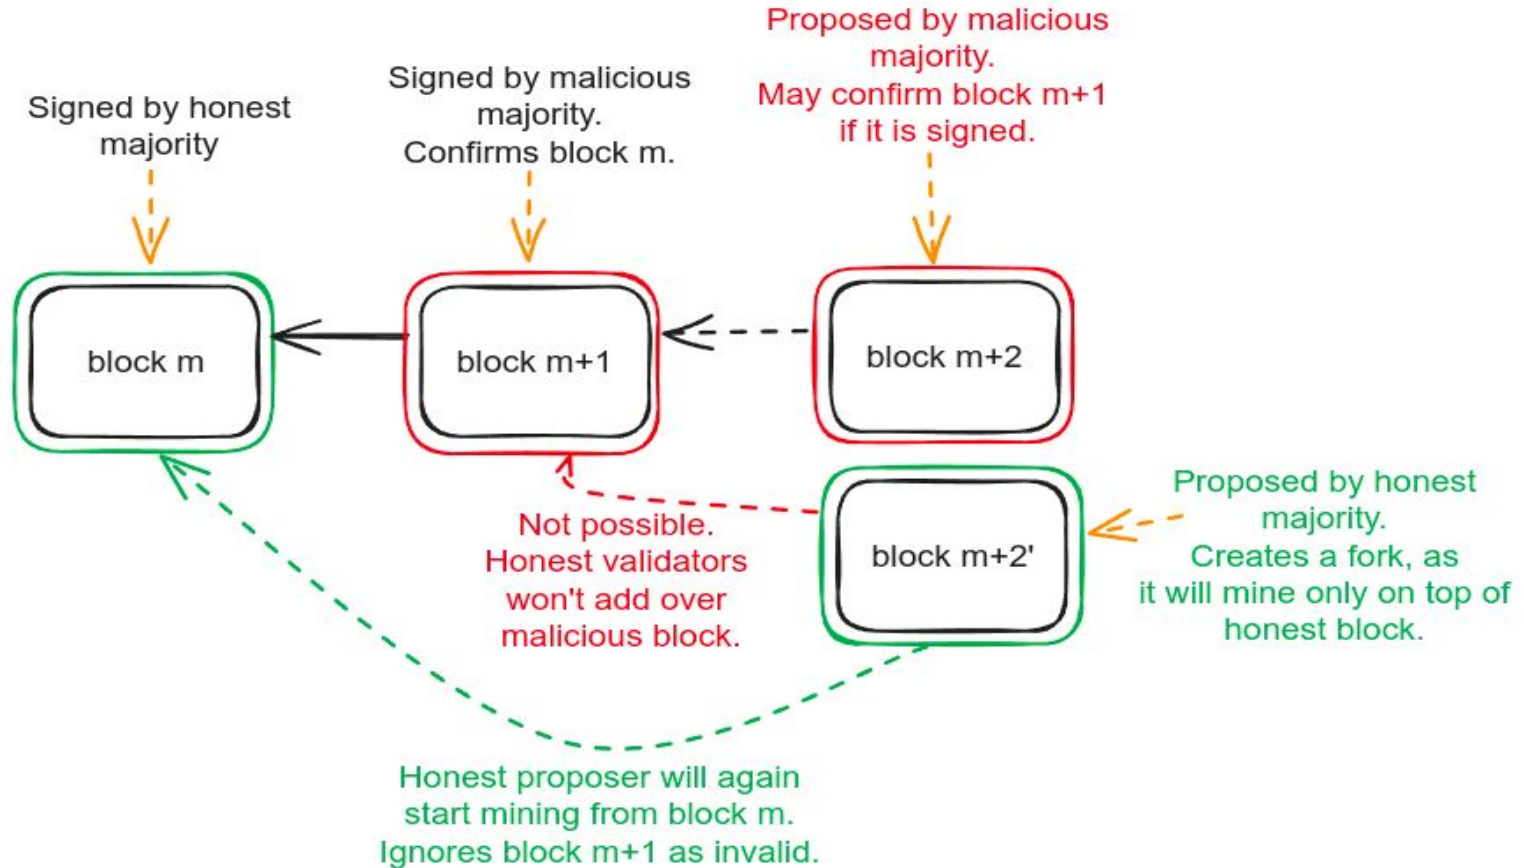
# Security



# Security




# Security



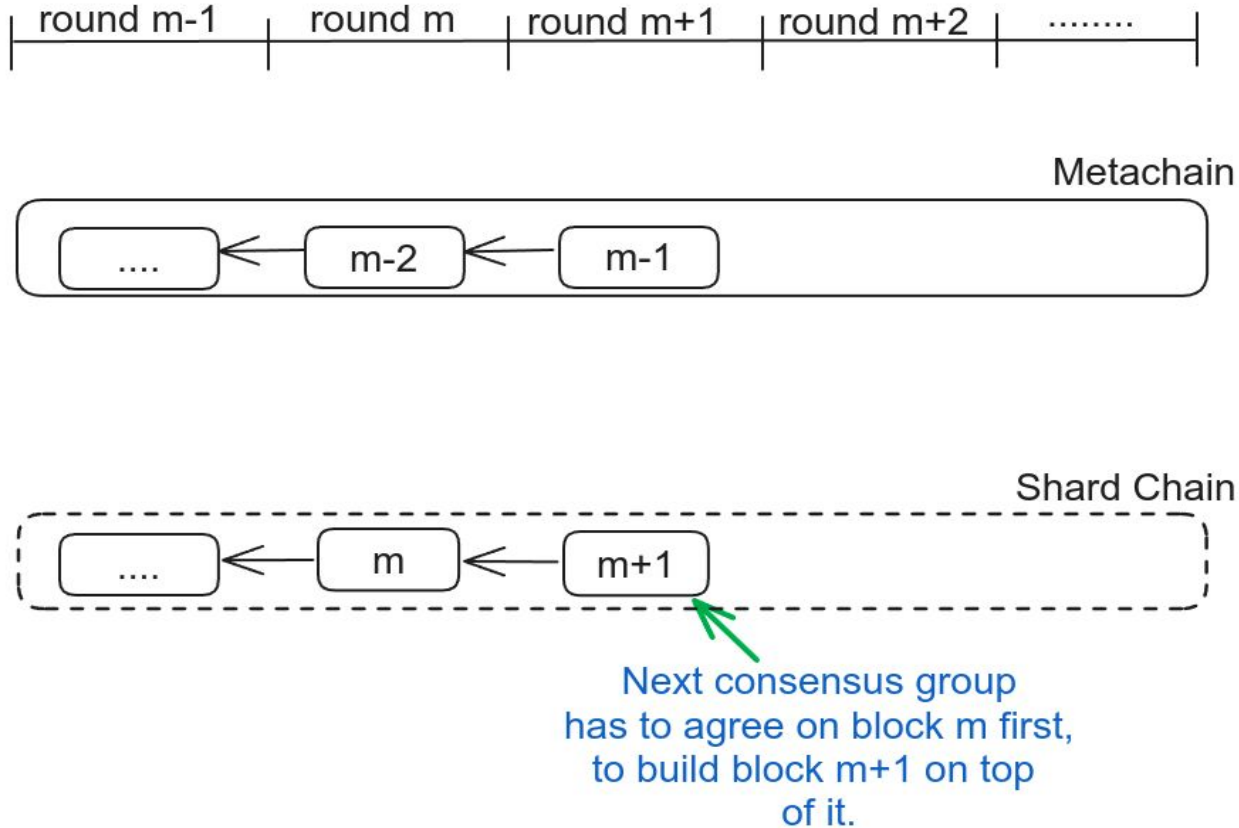
# Intra-Shard Transaction

- An Intra-Shard transaction can be called **processed** when it is included in a block and added to the respective chain of the shard associated with its *source address*.
- An Intra-Shard transaction can be called **confirmed** when the **header of the block** where the transaction is present is included in a block **added to the metachain**.

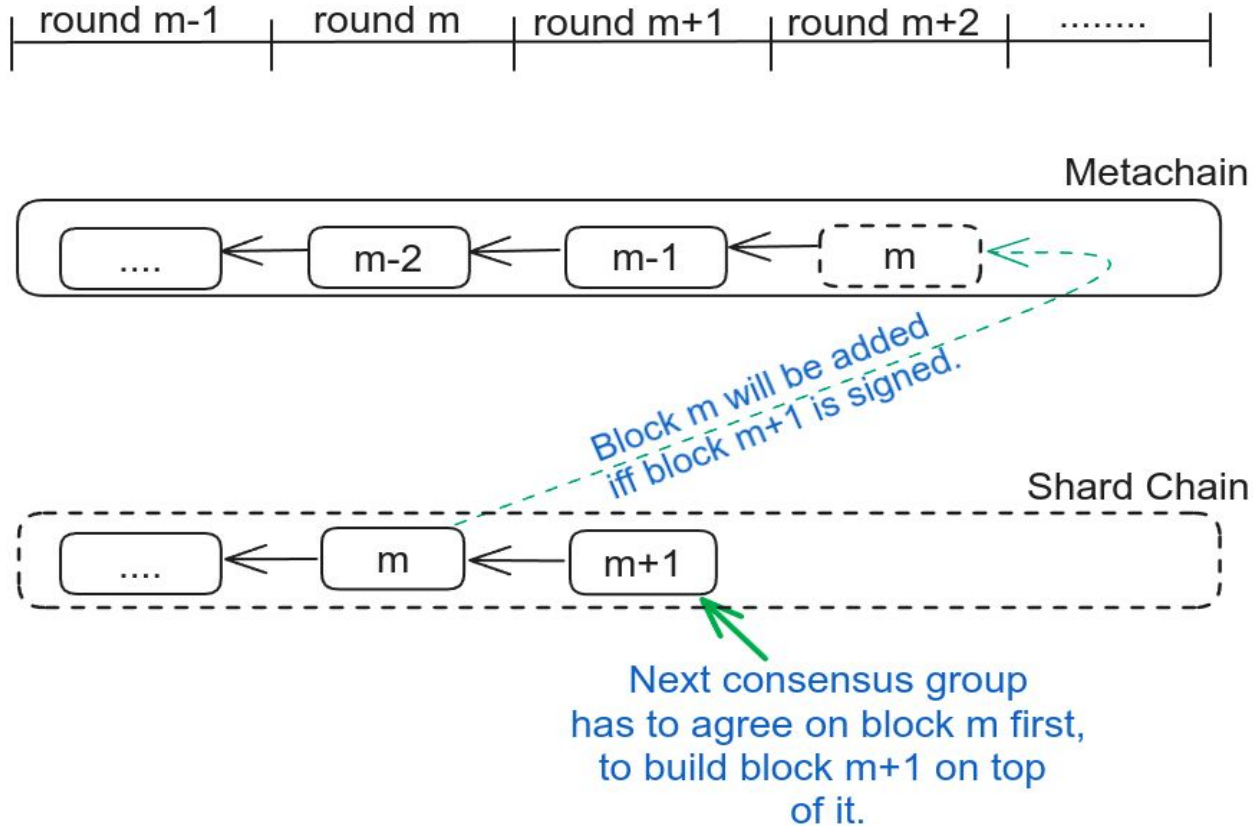
# Intra-Shard Transaction



# Intra-Shard Transaction




# Intra-Shard Transaction



# Intra-Shard Transaction

1 Round to Process

3 Rounds to Confirm






# Cross-Shard Transaction


- A Cross-Shard transaction can be called **processed** when it is **included in a block** and added to the respective chain of the **shard** associated with its *destination address*.
- A Cross-Shard transaction can be called **confirmed** when the **header of the cross-shard block** where the transaction is present is included in a block **added to the metachain**.

# Cross-Shard Transaction



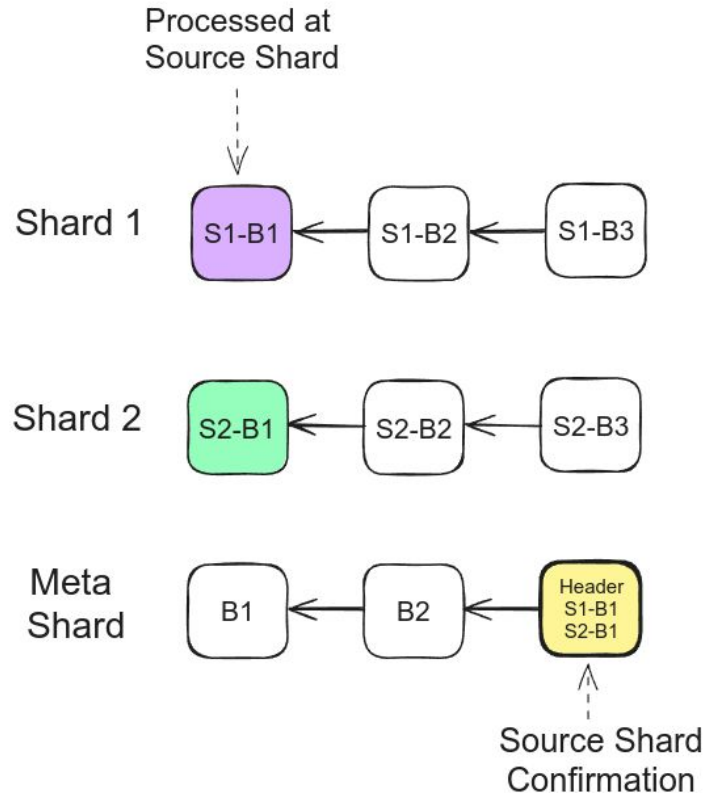
$S_i-B_j$  represents the  
 $j$ 'th block of shard " $i$ ".

# Cross-Shard Transaction



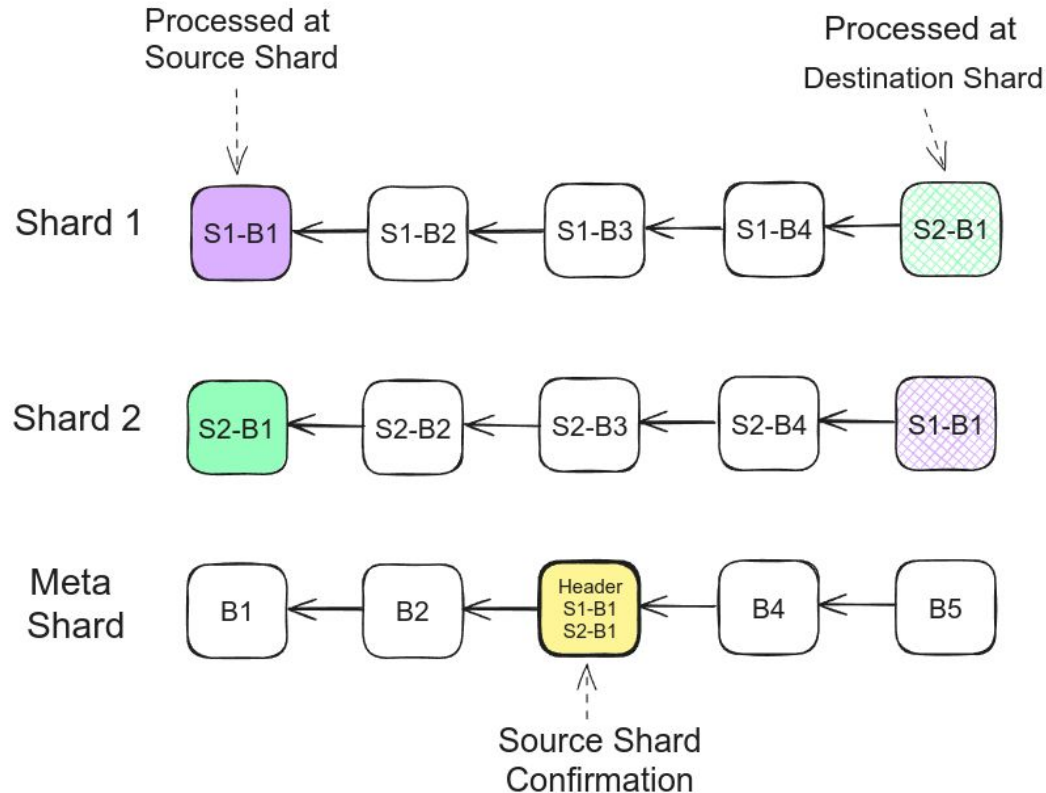
$S_i-B_j$  represents the  $j$ 'th block of shard " $i$ ".

# Cross-Shard Transaction



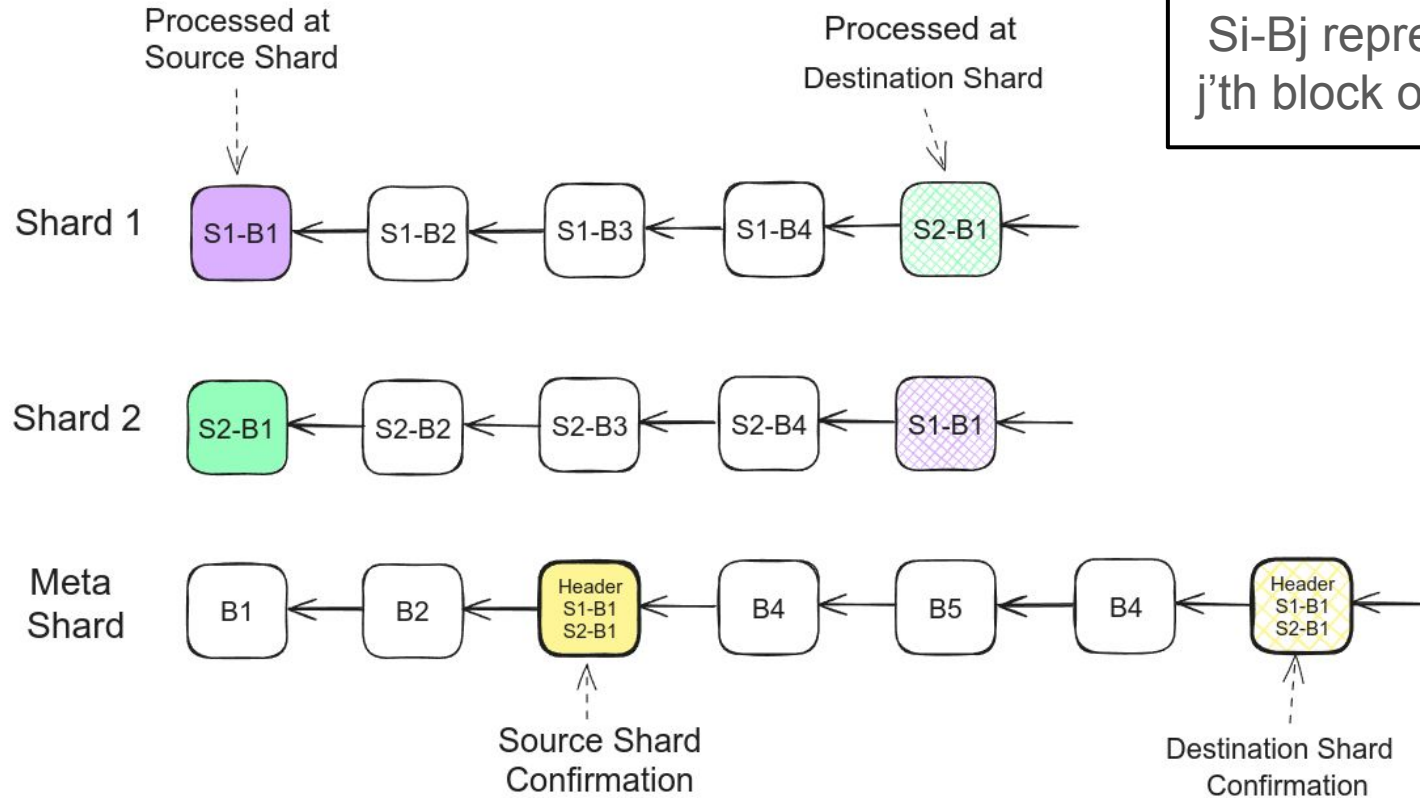
$S_i-B_j$  represents the  $j$ 'th block of shard " $i$ ".

# Cross-Shard Transaction

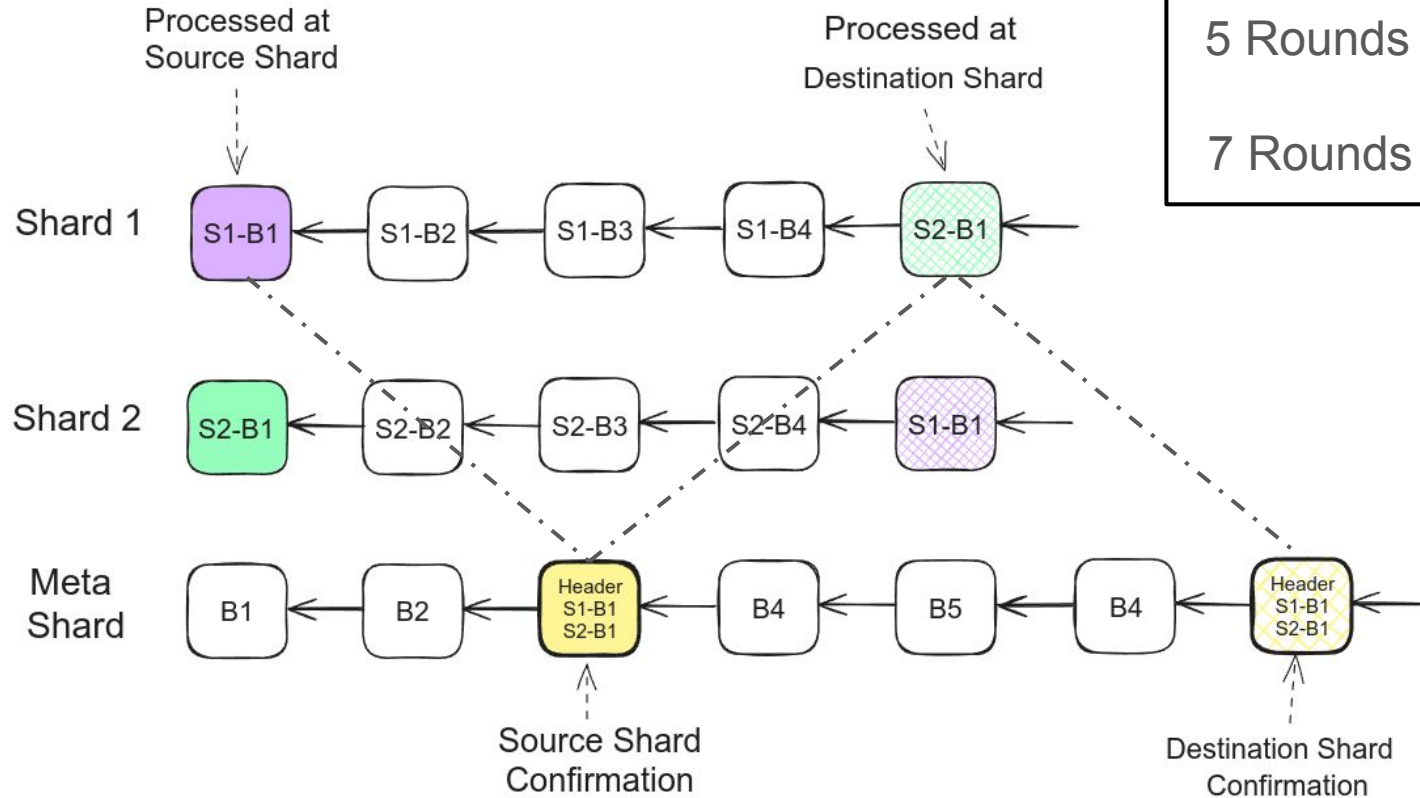


$S_i-B_j$  represents the  $j$ 'th block of shard " $i$ ".

# Cross-Shard Transaction



# Cross-Shard Transaction




# Experiment

- Simulated their local testnet with different configuration of shards. Integrated their Proxy Node and Go & Python SDK with the local testnet.
- Made a Load Testing repository in Python for interacting with the local testnet for creating accounts, loading various types of transaction into the network and functions to analyse the network status.
- Made changes in the local testnet code to work along with Load Testing script.
- Every experiment is run for 10 epochs of 100 rounds with different loads.
- Finally analysed the performance of the network and compared it with other popular models.



# Load - Intra-Shard Transactions

Loading only intra-shard transactions to every shard such that each shard reaches maximum throughput at each round.



# Load - Intra-Shard Transactions (Observations)

- The throughput almost **increases linearly** as expected.
- Whenever the **network reaches the threshold, number of shard increases**, to increase the overall throughput.
- The **metachain may act as the bottleneck** for confirmation latency if it gets full with shard block headers, thus requiring more than 1 meta block to confirm.
- But this is **not practical** in real time. When the number of shard increases, there can be **huge number of cross shard transactions** initiated.

# Load - Cross-Shard Transactions

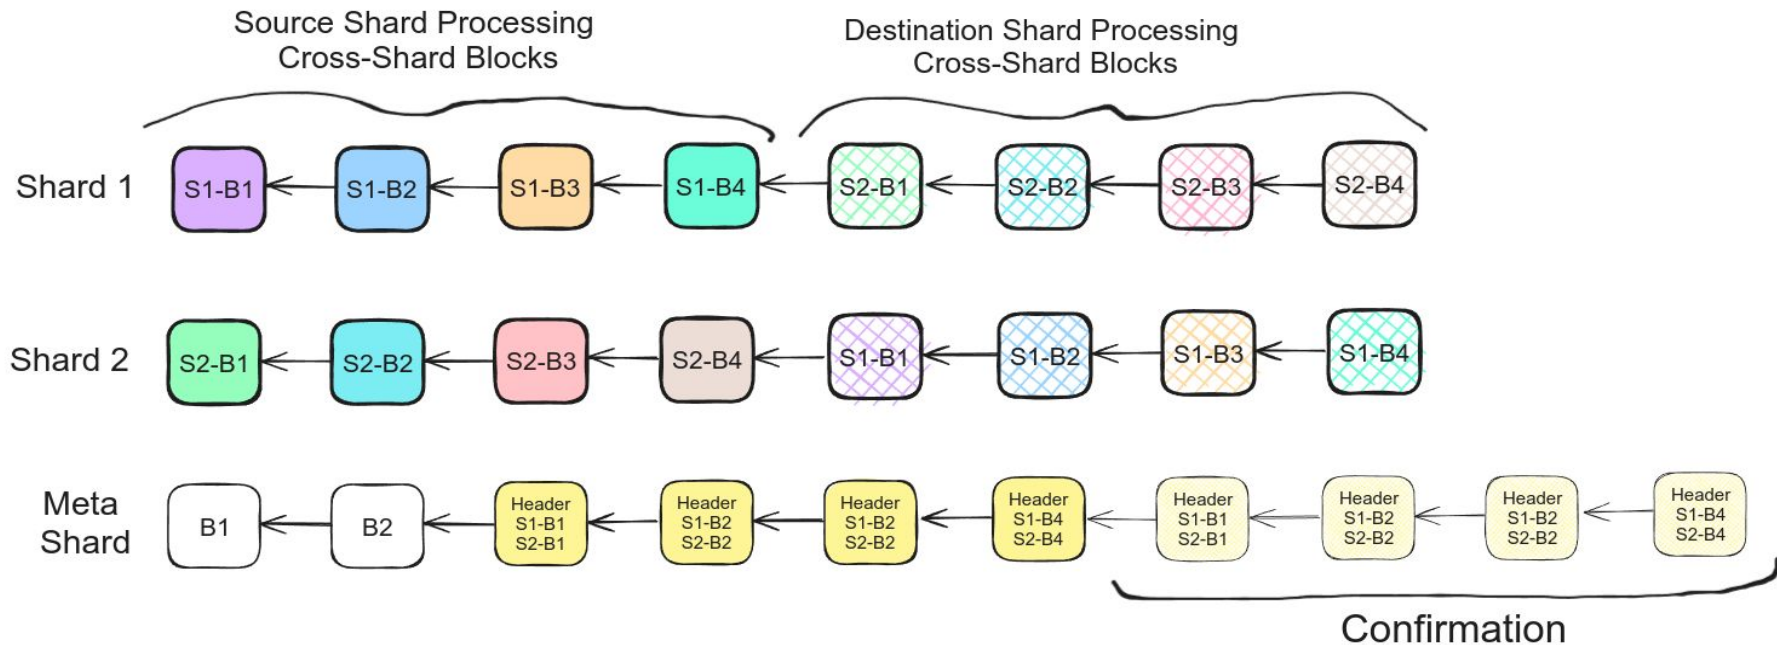
Cross-Shard transactions need 5 rounds for processing.

The transaction can be said processed only when it gets added to the chain of the destination shard.

# Load - Cross-Shard Transactions

Cross-Shard transactions need 5 rounds for processing.

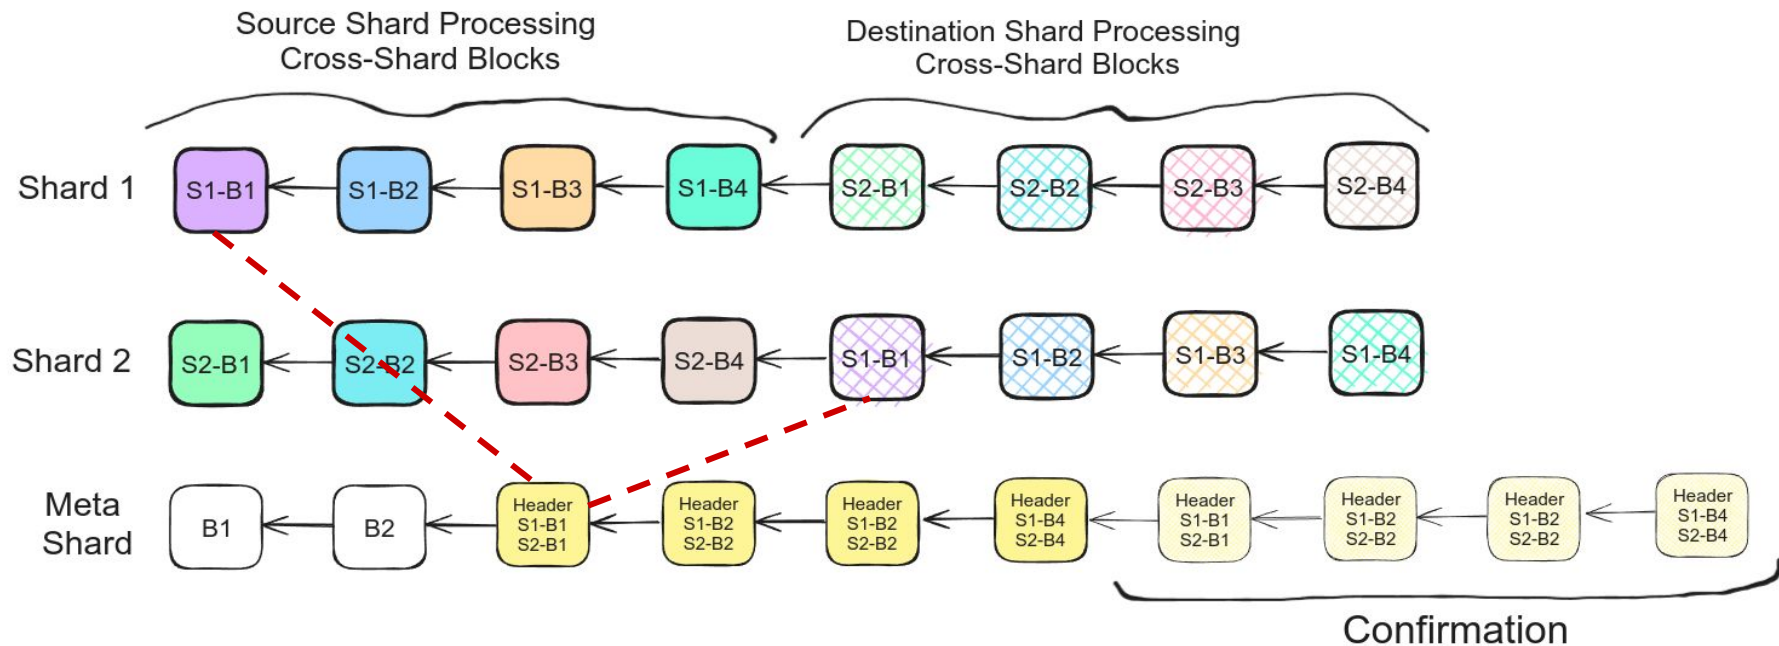
The transaction can be said processed only when it gets added to the chain of the destination shard.



# Load - Cross-Shard Transactions

Cross-Shard transactions need 5 rounds for processing.

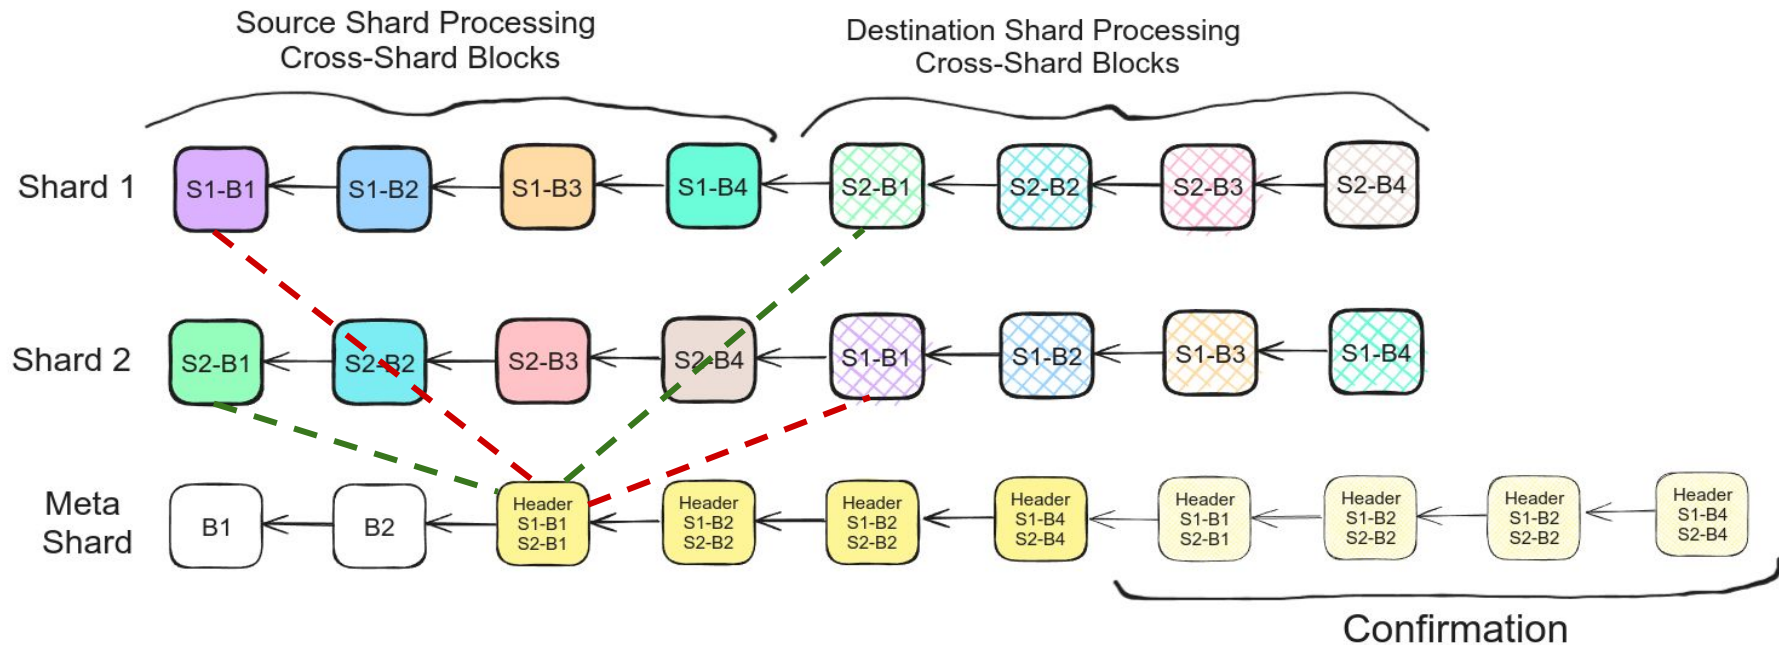
The transaction can be said processed only when it gets added to the chain of the destination shard.



# Load - Cross-Shard Transactions

Cross-Shard transactions need 5 rounds for processing.

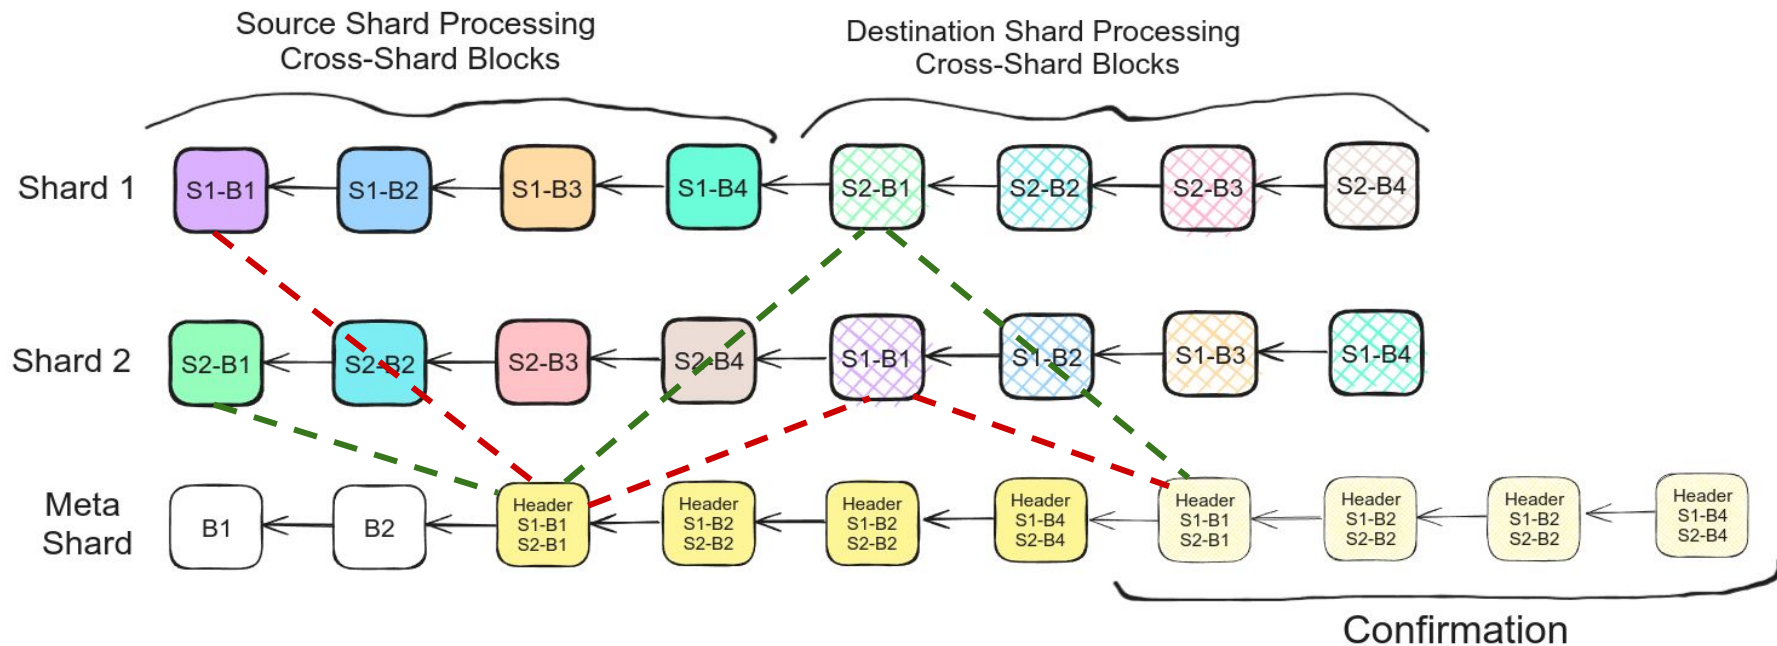
The transaction can be said processed only when it gets added to the chain of the destination shard.




# Load - Cross-Shard Transactions

Cross-Shard transactions need 5 rounds for processing.

The transaction can be said processed only when it gets added to the chain of the destination shard.



# Load - Cross-Shard Transactions






# Load - Cross-Shard Transactions (Observations)

- Every transaction's **destination account has been chosen uniformly** from all the accounts in the network.
- The **cross-shard gives exactly half the throughput** achieved in Intra-Shard when every block is filled up i.e. (Peak Throughput).
- The peak throughput is only achievable when a shard is receiving maximum transactions per second from other shard.  
Otherwise, **increase in shard may lead to more number of cross-shard transactions and thus a dip in the throughput.**
- There can be some cases where one shard gets overwhelmed with cross-shard transactions from other shards.

# Load - Mixed-Shard Transactions




# Observations

- Every cross transaction's destination account has been chosen uniformly from all the accounts in the network.
- With **increase in number of shards**, the **throughput increases** linearly even in the case of Mixed transactions.
- The same assumption is made here that every block is receiving full transactions, to get measurement of the peak throughput.


# Load - Cross-Shard Transactions (OvA)

Send Cross-Shard transaction to only one specific Shard.

Shards

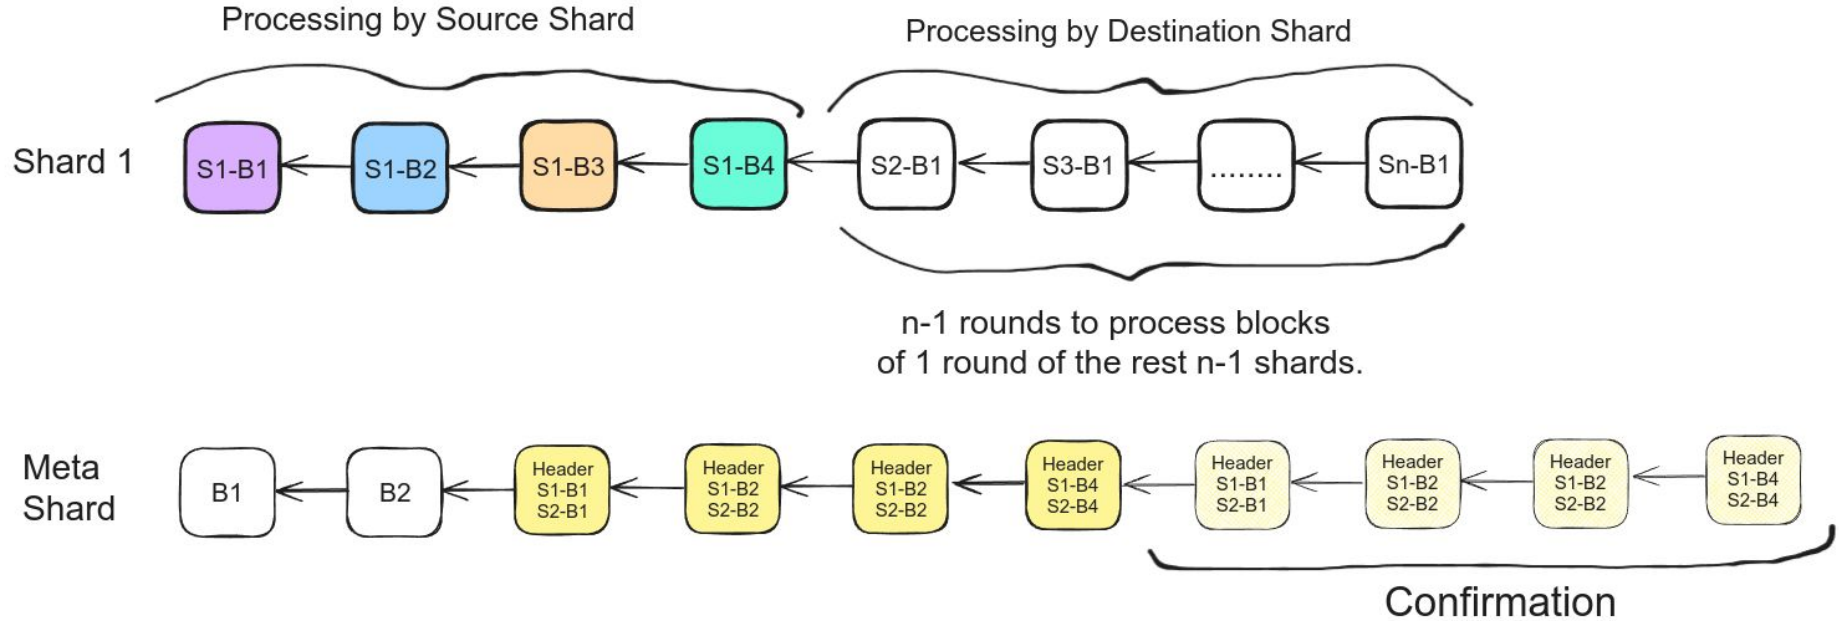


Before




Now  
One Vs. All (OvA)

# Load - Cross-Shard Transactions (OvA)



# Load - Cross-Shard Transactions (OvA)



# Observations

- With a timeline of 5 epochs where the **source shard can create 4 blocks at the first 4 rounds**, the **rest of the rounds it spent finalizing the cross-shard blocks from other transactions**. Thus, even if we increase the size of the shard, the throughput remains const approx **3500 TPS**.
- Shard Re-Adjustment or any other method must be used to distribute the cross shard transactions to more shards.


# Increase in Percentage of Cross-Shard with Increase in Shard

As we increase the number of shards,  
what is the percentage increase in cross-shard txn?

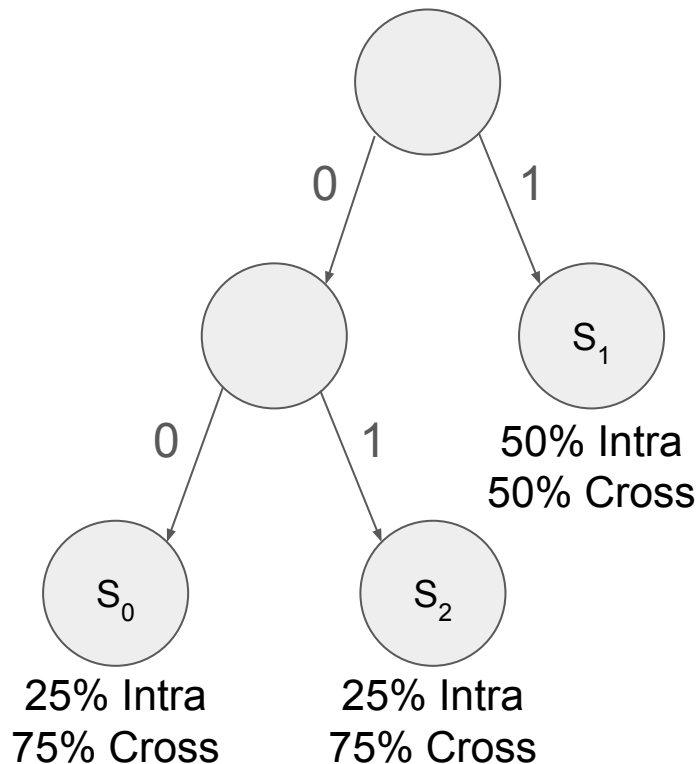

We experiment for a fixed period of 5 epochs and 1000 accounts.  
The accounts for transactions are uniformly picked up from all the shards.



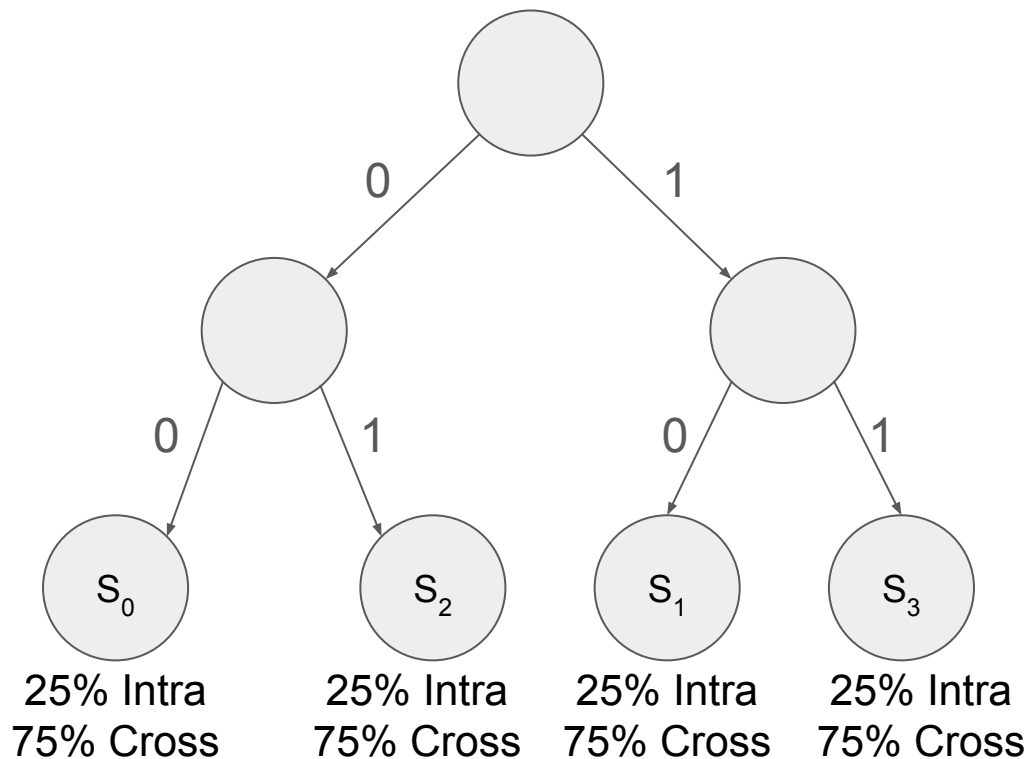
# Increase in Percentage of Cross-Shard with Increase in Shard



# Increase in Percentage of Cross-Shard with Increase in Shard

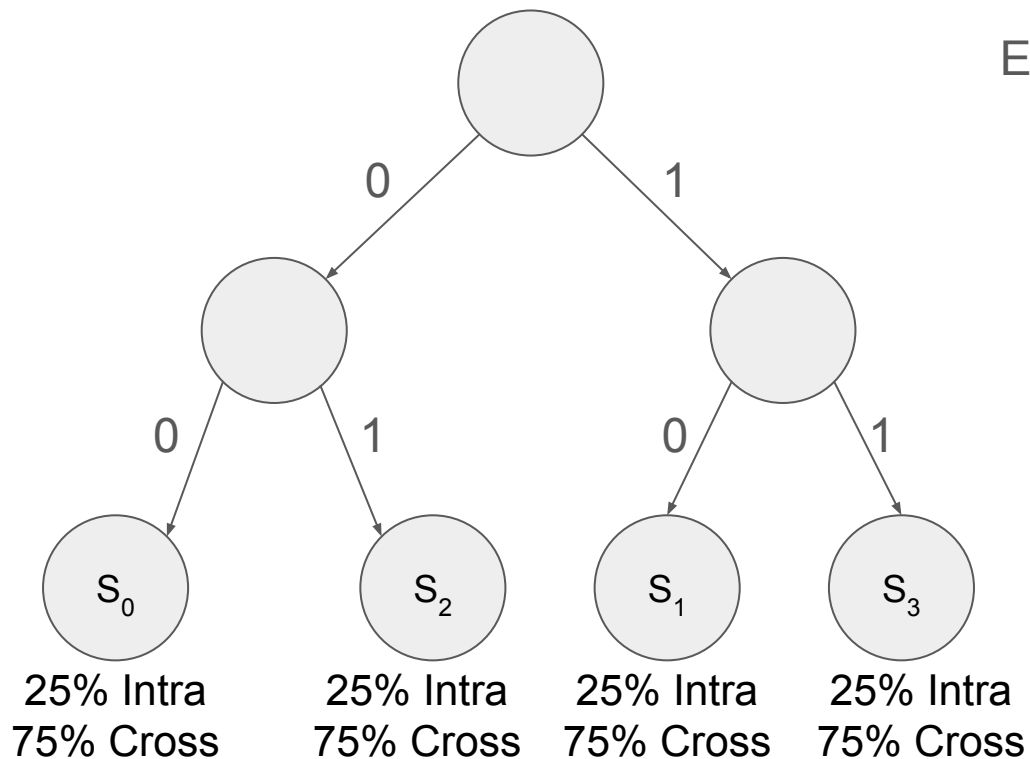


# Increase in Percentage of Cross-Shard with Increase in Shard

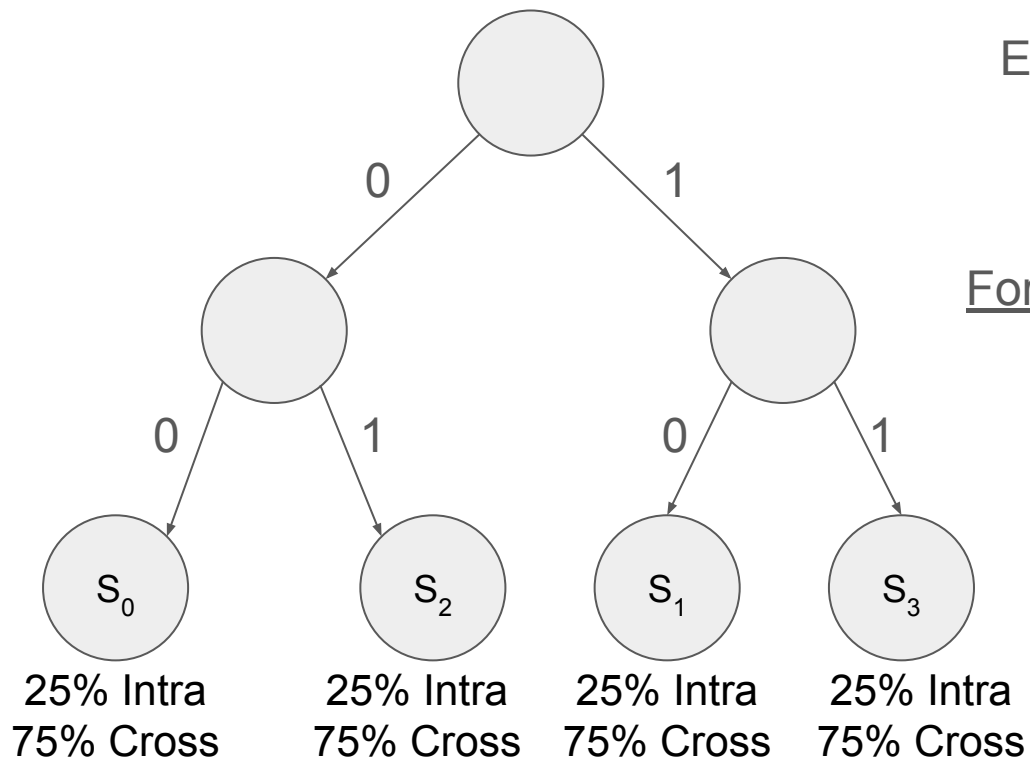


# Increase in Percentage of Cross-Shard with Increase in Shard

Each Block in a Round (20k Txns)  
Intra-Shard: 5000 Txns  
Cross-Shard: 15000 Txns



# Increase in Percentage of Cross-Shard with Increase in Shard



Each Block in a Round (20k Txns)  
Intra-Shard: 5000 Txns  
Cross-Shard: 15000 Txns

For consecutive 8 Rounds of a Shard

5000 Intra-Shard
15000 Source Cross-Shard

First 4  
Rounds

5000 Intra-Shard
15000 Destination Cross-Shard

Next 4  
Rounds


# Increase in Percentage of Cross-Shard with Increase in Shard

Each Block in a Round (20k Txns)

Intra-Shard: 5000 Txns

Cross-Shard: 15000 Txns

For consecutive 8 Rounds of a Shard



First 4  
Rounds



Next 4  
Rounds

For 1 Shard:

#Intra =  $5000 \times 8 = 40000$

#Cross =  $15000 \times 4 = 60000$

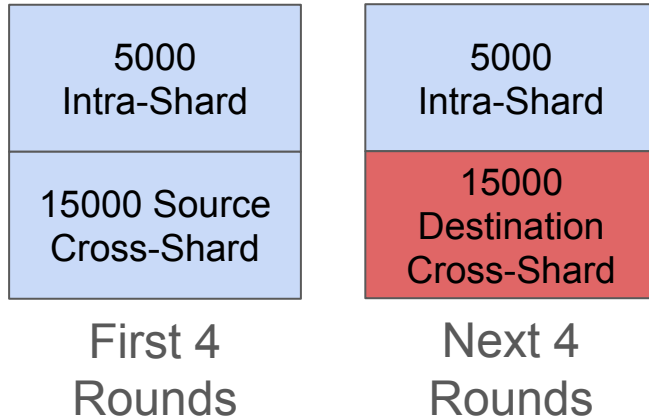
# Increase in Percentage of Cross-Shard with Increase in Shard

Each Block in a Round (20k Txns)

Intra-Shard: 5000 Txns

Cross-Shard: 15000 Txns

For consecutive 8 Rounds of a Shard



For 1 Shard:

$$\# \text{Intra} = 5000 \times 8 = 40000$$

$$\# \text{Cross} = 15000 \times 4 = 60000$$

For 4 Shards:

$$\# \text{Intra} = 40000 \times 4 = 160000$$

$$\# \text{Cross} = 60000 \times 4 = 240000$$

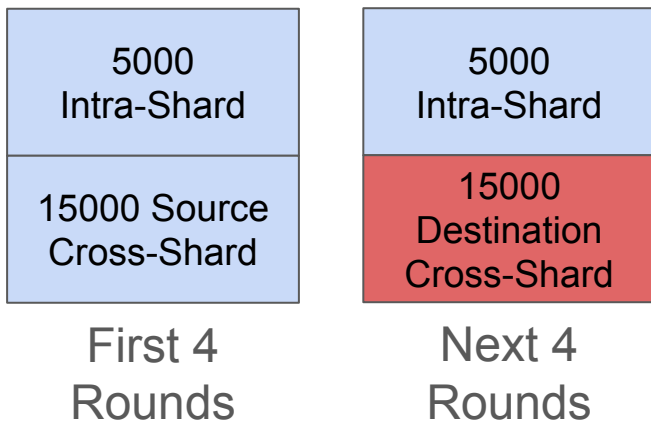
# Increase in Percentage of Cross-Shard with Increase in Shard

Each Block in a Round (20k Txns)

Intra-Shard: 5000 Txns

Cross-Shard: 15000 Txns

For consecutive 8 Rounds of a Shard



For 1 Shard:

#Intra =  $5000 \times 8 = 40000$

#Cross =  $15000 \times 4 = 60000$

For 4 Shards:

#Intra =  $40000 \times 4 = 160000$

#Cross =  $60000 \times 4 = 240000$


#Percentage of Cross-Shard = 60%  
(Same as received in Experiment)



# Increase in Percentage of Cross-Shard with Increase in Shard




# Minimising Cross-Shard by Opening Additional Accounts




Lots of Transactions  
between these two  
accounts.

# Minimising Cross-Shard by Opening Additional Accounts




Lots of Transactions  
between these two  
accounts.




Create new wallet for  
one of the account in  
the other shard.

# Minimising Cross-Shard by Opening Additional Accounts



Lots of Transactions  
between these two  
accounts.



Create new wallet for  
one of the account in  
the other shard.

Transaction can be  
processed 2x speed

# Minimising Cross-Shard by Opening Additional Accounts



Lots of Transactions  
between these two  
accounts.

Create new wallet for  
one of the account in  
the other shard.

Transaction can be  
processed 2x speed

Overheads

# Minimising Cross-Shard by Opening Additional Accounts

- Accounts with **more than 3 transactions between** them in a round are optimized with this technique.
- **With 4 shards, and 1000 accounts**, and blocks full of transactions, the percentage of Cross-Shard could be **reduced to around 48% from 60%**.

# Conclusion

- MultiVersX **successfully implemented sharding** and currently running with 3 shards on the main network.
- With almost **3.4M user accounts** in the mainnet, it comes out as one of the most successful sharded system currently.
- Their VM supports writing of **Smart Contracts with Rust** as it runs on Web Assembly.
- **Bottlenecks** created by Cross-Shard, Meta-Chain and Security Issues need to be overcome to make sharding a goto idea for scalability.
- With this R&D, **learnt about practical implementation of sharding** along with how to do analysis of sharded system.

# References

- [1] Multiversx whitepaper.
- [2] Miguel Castro and Barbara Liskov. Practical byzantine fault tolerance. In OsDI, volume 99, 1999.
- [3] Sunny King and Scott Nadal. Ppcoin: Peer-to-peer crypto-currency with proof-of-stake. self-published paper, 19(1), August 2012.
- [4] Paradigm Research. Elrond: High-throughput public blockchain with adaptive state sharding, 2019.
- [5] A. Secure. The zilliqa project: A secure, scalable blockchain platform. 2018.
- [6] Alex Skidanov and Illia Polosukhin. Nightshade: Near protocol sharding design. 2019.
- [7] Dmitry Tanana. Avalanche blockchain protocol for distributed computing security. In 2019 IEEE International Black Sea Conference on Communications and Networking (BlackSeaCom). IEEE, 2019.
- [8] Gavin Wood. Polkadot: Vision for a heterogeneous multi-chain framework. White paper, 21(2327):4662, 2016.